Choose 5-6 pictures of generated images to show how training progresses (for example - epoch 0, 100, 200, 300, 400, 500)

Epoch 0

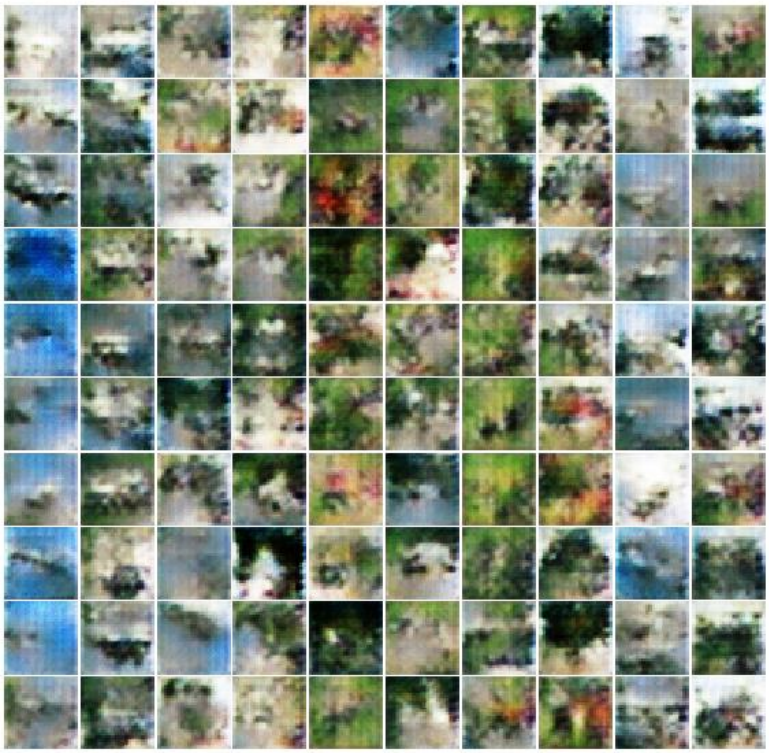

Epoch 10

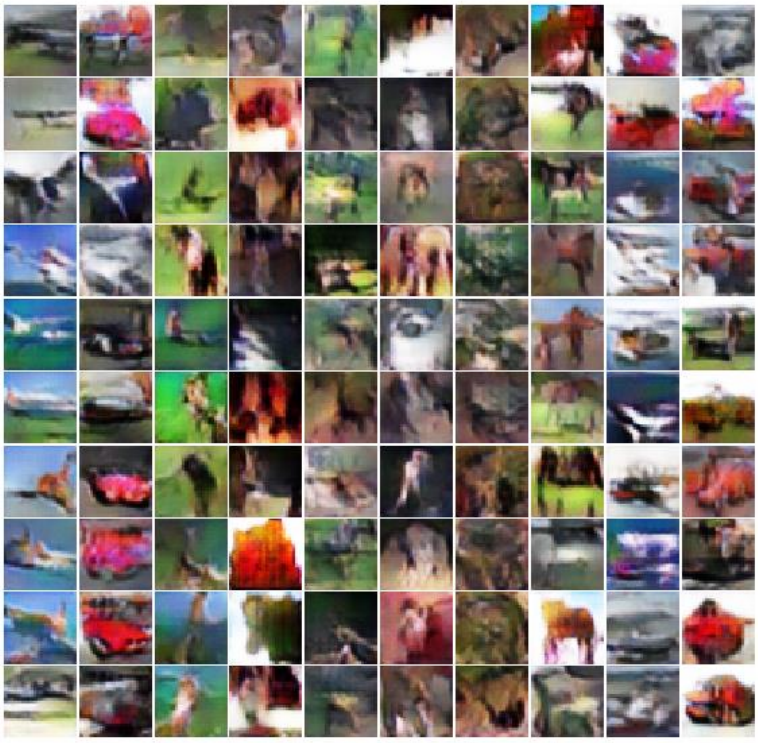

Epoch 50

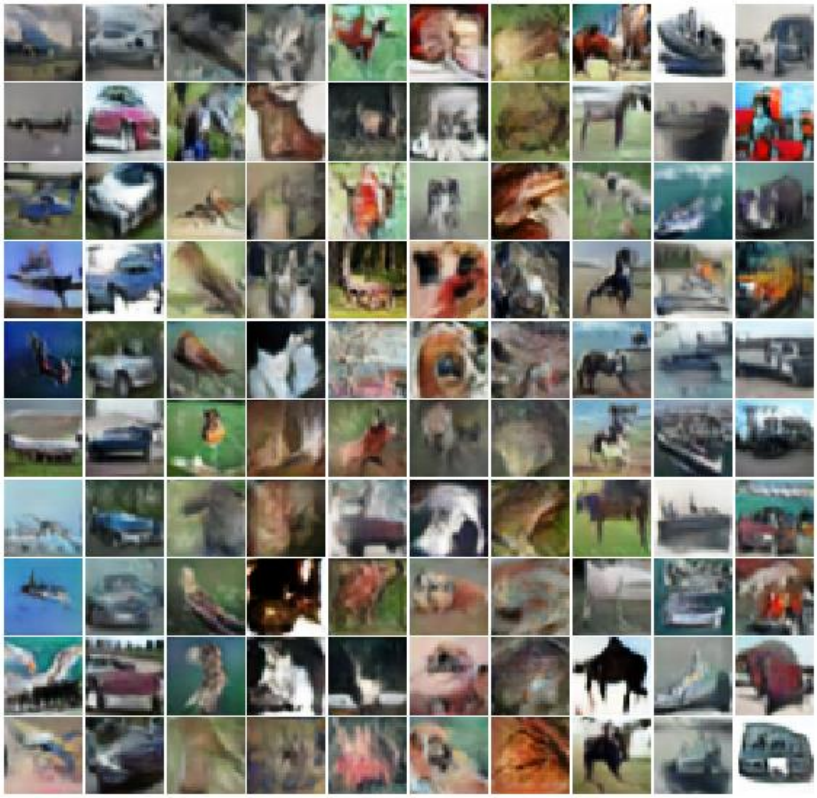

Epoch 100

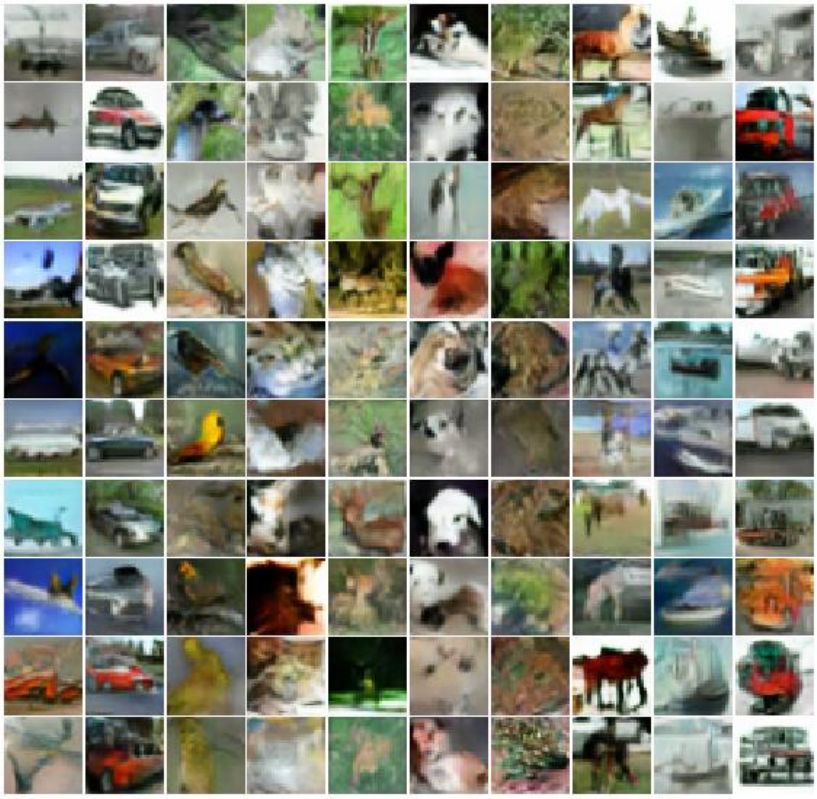

Epoch 200

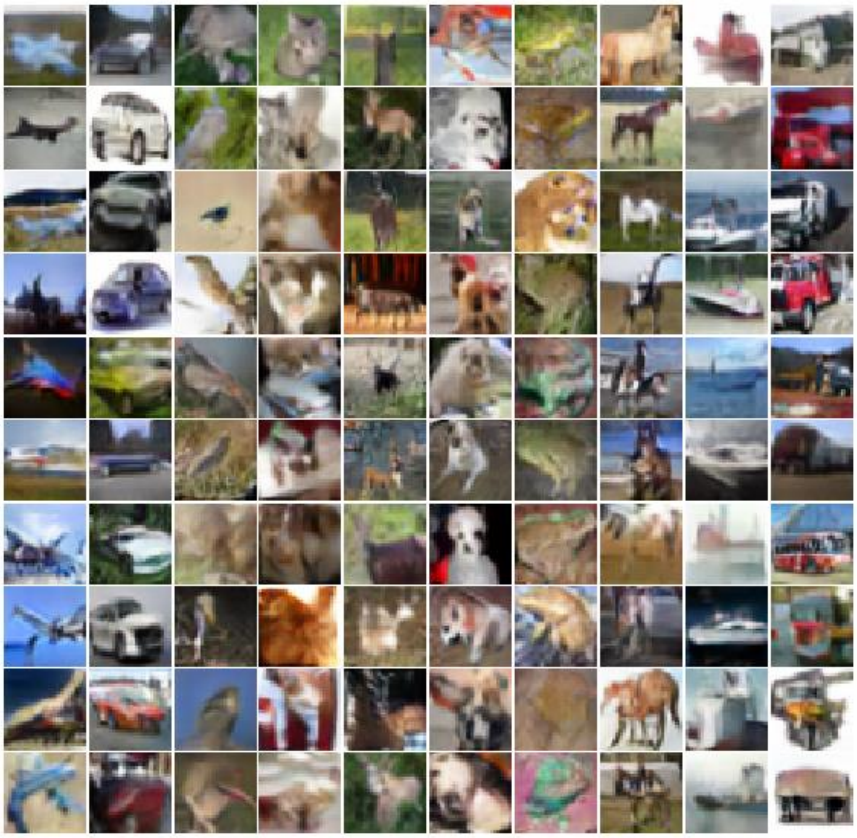

Epoch 400

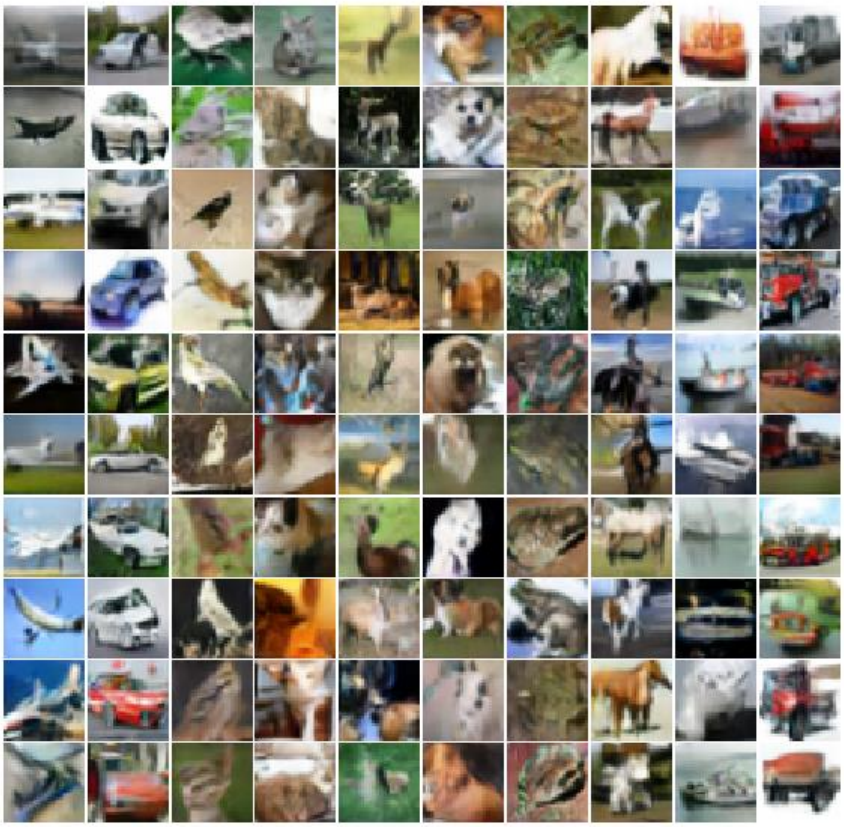

Submit a batch of real images, a batch of the gradients from an alternate class for these images, and the modified images the discriminator incorrectly classifies

Real Images

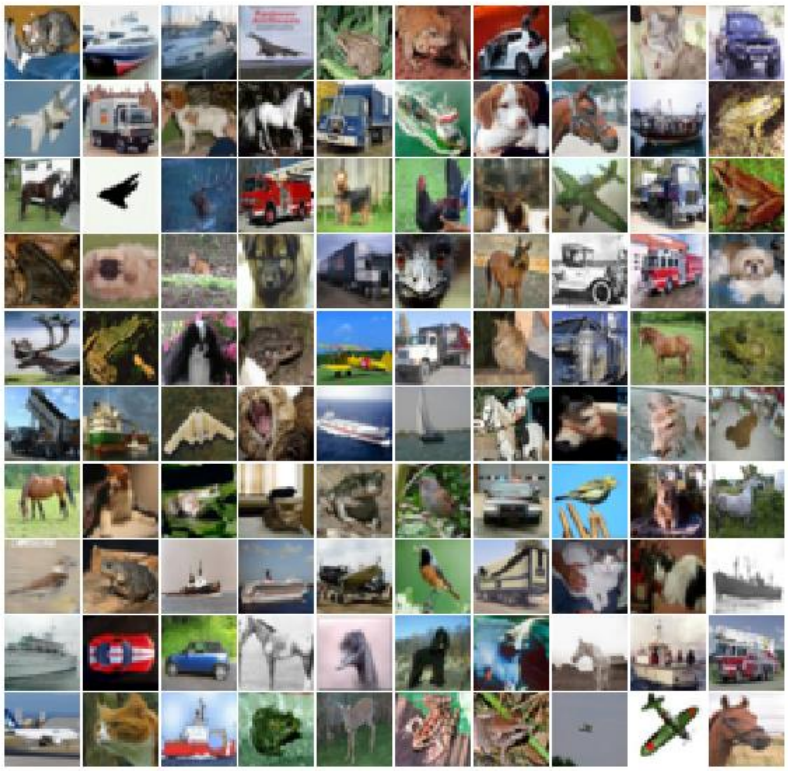

Gradients

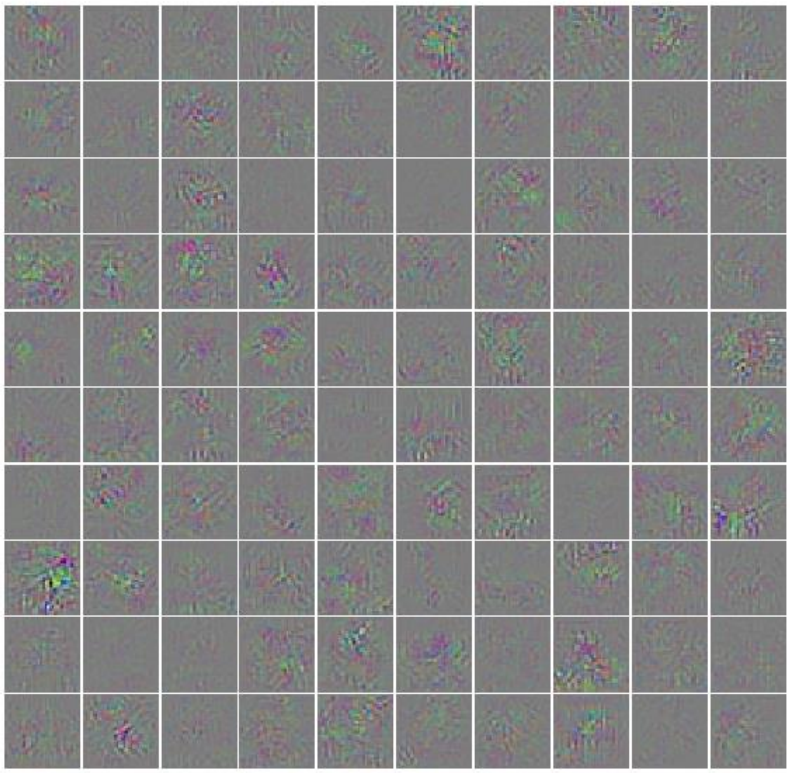

Modified Images

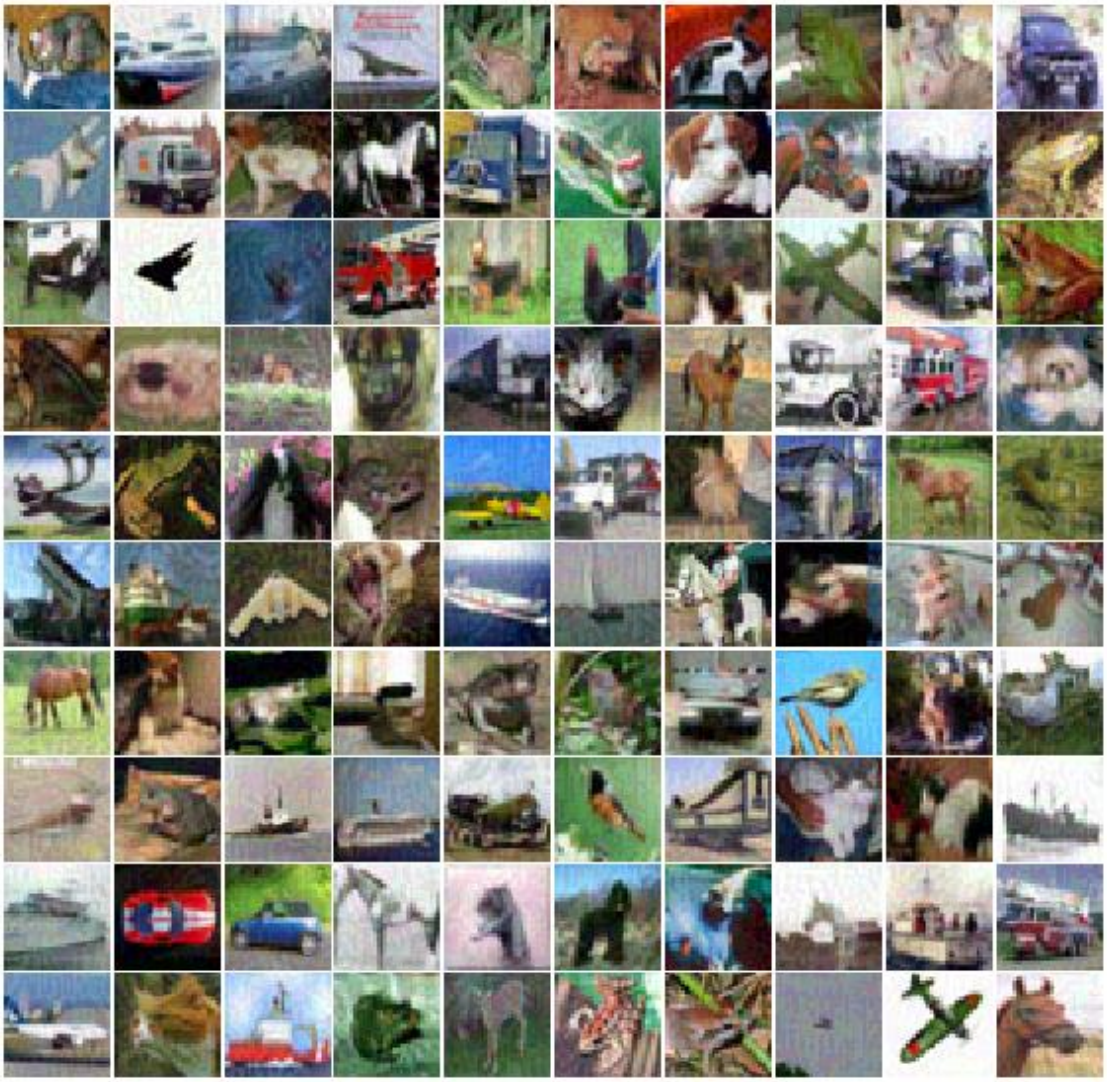

Submit synthetic images maximizing the class output. One for the discriminator trained without the generator and one for the discriminator trained with the generator.

Discriminator trained without the generator

Discriminator trained with the generator

Submit synthetic images maximizing a particular layer of features. Do this for at least two different layers (for example - layer 4 and layer 8.)

Discriminator trained without the generator
Layer 4:

Layer 8:

Discriminator trained with the generator

Layer 4:

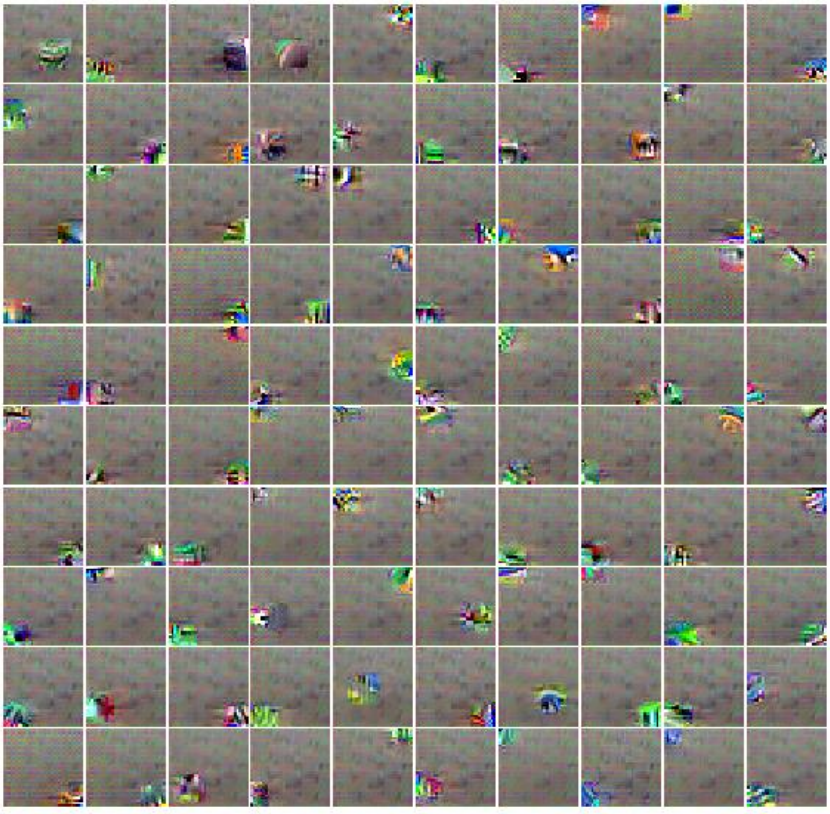

Layer 8:

Report your test accuracy for the two discriminators.

Batch Size = 64

Number of Features = 128

Discriminator trained without the generator
Number of Epochs: 100
Test Accuracy: 87%
Train accuracy: 98%

Discriminator trained with the generator
Number of Epochs: 400
Test Accuracy: 83%
Train accuracy: 92%

Part 2: Perturb Real Images

Accuracy on Real Images: 93%

Accuracy on Altered Images: 15%